

AGENTIC AI CONSULTING BUSINESS FRAMEWORK

From Discovery to Architecture to Build to Deploy to Monitor to Monetise

A complete operating system for Agentic AI Solutions Consultants
with Product Manager and Product Builder mindset

Multi-Agent Architecture | Cloud Infra Costs | LLM Pricing
Observability ROI | Proposal Templates | Monetisation Ladder

End to End flow of Agentic AI




Table of Contents

1. Your Consulting Identity and Positioning
2. The Consulting Flywheel Mental Model
3. Ideal Client Profile
4. Approach Framework: Discovery to Pitch to Win
5. AI Opportunity Scoring Matrix
6. Solution Types and Pricing Ladder
7. Project Lifecycle (8 Phases)
8. Client Expectation Management
9. Solution Architectures (4 Types)
10. LLM Model Cost Comparison (2025)
11. Cloud Infrastructure and Cost Structure
12. Observability and Cost Management ROI
13. Proposal and Delivery Structure
14. Monetisation Model and Retainer Strategy
15. Your 90-Day Business Launch Plan
16. Final Mental Model and Differentiation

1. Your Consulting Identity and Positioning

Core Positioning Statement

I help companies design, build, and deploy secure, observable, cost-efficient Agentic AI and multi-agent systems that automate business workflows, reduce manual effort, and produce measurable ROI.

What You Are NOT Selling

- Generic AI chatbots or one-off demos
- One-size-fits-all AI tools
- Expensive tech nobody can maintain

What You ARE Selling

Value Delivered	How
Saved Hours	Automation of repetitive knowledge workflows
Cost Reduction	30 to 85 percent LLM cost savings through routing, caching, prompt optimisation
Operational Clarity	Observable AI with dashboards, traces, and cost alerts
Decision Support	Agents that augment, not replace, human judgment
Measurable ROI	Business metrics tied directly to AI performance

Your Role Map

Role	Responsibility
Business Discovery	Identify the right workflows and build the business case
Product Manager	Define scope, success metrics, risk, and user stories
Solution Architect	Design LLM, agent, cloud, and data architecture
AI Builder	Implement RAG, agents, tools, APIs, and pipelines
Observability Expert	Set up tracing, cost dashboards, alerting, and evaluation
Delivery Manager	Manage milestones, expectations, and handover documentation
Retainer Partner	Ongoing optimisation, improvement, and support

2. The Consulting Flywheel Mental Model

Every engagement follows this loop. Internalise it as your operating system.

Flywheel Step	What You Do
Business Pain	Listen for slow, manual, expensive knowledge workflows
Workflow Discovery	Map exactly what happens, who does it, how long it takes
Agent Opportunity	Identify where an AI agent creates value, not just automation
ROI Estimation	Quantify savings: hours x rate + incidents + waste
MVP Proposal	Smallest solution that proves value to the business
Architecture Design	LLM, RAG, tools, cloud, observability, security
Build and Integrate	Implement, test, and connect to real business systems
Evaluate and Monitor	Measure quality, cost, latency, hallucination, user feedback
Cost Optimise	Prompt tuning, caching, model routing, alert thresholds
Retainer and Expand	Monthly optimisation and expand to next workflow

3. Ideal Client Profile

Company Type	Primary AI Opportunity
SaaS Companies	Customer support, onboarding, internal ops automation
IT Services	Ticket triage, code review, documentation agents
Finance / FinTech	Report generation, compliance review, document analysis
Healthcare Admin	Intake automation, claims support, patient summarisation
Legal Firms	Contract review, legal research, case summarisation
Manufacturing	SOP assistants, maintenance agents, procurement automation
HR / Recruitment	Resume screening, interview assistant, policy bot
Education / Training	Tutor agents, course assistants, content generation
Enterprise Support	Knowledge-base RAG agent, escalation assistant
Sales Organisations	CRM agent, lead research, proposal generators

Best First Qualifying Signals

- Employees spend hours searching documents, PDFs, or internal wikis
- Support or ops teams answer the same questions repeatedly
- Reports are generated manually by copying data between tools
- Ticket or request backlog overwhelms the team
- Decisions are delayed because information cannot be found quickly
- Onboarding takes weeks because knowledge is not centralised

4. Approach Framework: Discovery to Pitch to Win

Step 1 -- Cold Outreach Template

Cold Outreach Message

Hi [Name],

I help companies build production-ready AI agents that automate repetitive knowledge workflows: support, internal search, document analysis, and tool-based operations.

A common opportunity: an AI agent that answers from company documents, calls internal APIs, creates structured outputs, escalates to humans when needed, and tracks cost and accuracy.

Would you be open to a 20-minute AI Workflow Audit to identify 1 to 2 areas where an agent could save time or reduce cost?

Step 2 -- Discovery Call Questions

Business Process Questions

- Which workflows are repetitive and time-consuming?
- Which teams spend the most time searching for information?
- Where do employees copy and paste between systems?
- What tasks are high-volume but low-complexity?
- Which processes require approvals or escalation?

Cost and ROI Questions

- How many people work on this process and how many hours per week?
- What is the approximate hourly cost of those people?
- How many requests, tickets, or documents happen per month?
- What happens when this process is delayed?
- What is the cost of mistakes or rework?

AI Readiness Questions

- Where is the data stored? Structured or unstructured?
- Which tools or APIs need integration?
- Are there compliance or privacy (PII / GDPR) requirements?
- Does output need human approval before action?
- Is cloud access available, or must it run on-premises?

Success Definition Questions

- What metric should improve -- cost, time, accuracy, or volume?
- What is the acceptable accuracy level for production use?
- When should the AI escalate to a human?
- What does a successful 90-day outcome look like?

5. AI Opportunity Scoring Matrix

Score each use case 1 to 5. Prioritise use cases with the highest total score.

Criteria	Score	What to Assess
Repetition Frequency	1-5	How often does this workflow happen?
Manual Effort	1-5	How much human time does it consume?
Data Availability	1-5	Are documents, APIs, or databases accessible?
Business Impact	1-5	Does it save cost or improve revenue?
Risk Level	1-5	5 = very low risk to automate
Integration Complexity	1-5	5 = simple integrations needed
ROI Visibility	1-5	How easily can results be measured?

Prioritise Use Cases Where

Repetition >= 4 AND Data Available >= 4 AND Risk <= 2 AND ROI Visible >= 4
These are your safest, highest-value first engagements.

6. Solution Types and Pricing Ladder

Use this ladder to grow every client from audit to retainer. Never start with a complex system.

Offer	Price	Duration	Deliverables
Offer 1: AI Workflow Audit	\$500 - \$2,000	3-7 days	Workflow map, top 3 automation opportunities, ROI estimate, architecture recommendation, roadmap
Offer 2: Agentic AI MVP Build	\$3,000 - \$15,000	2-6 weeks	RAG over company docs, LLM answer generation, vector DB, API or UI, observability, evaluation, deployment
Offer 3: Production Agentic AI System	\$15,000 - \$75,000	6-16 weeks	Full cloud architecture, auth, DB integration, tool calling, human-in-the-loop, monitoring, guardrails, CI/CD
Offer 4: Multi-Agent System	\$30,000 - \$150,000+	8-20 weeks	Supervisor + specialist agents, cross-functional workflow, complex tool orchestration, full observability
Offer 5: Observability and Cost Retainer	\$2K-\$10K setup + \$500-\$5K/mo	2-4 weeks setup	LangFuse/LangSmith setup, token tracking, prompt cost analysis, model routing, caching, dashboards, monthly reports

7. Project Lifecycle -- 8 Phases

Phase	Name	Key Activities
Phase 1	Discover	Interview stakeholders, map workflow, identify data, estimate effort, define success metrics
Phase 2	Define	Write product spec: problem, users, inputs/outputs, tools, guardrails, human approval rules, non-goals
Phase 3	Design	Architecture: LLM, agent framework, vector DB, memory, cloud, observability, security, cost management
Phase 4	Build	Implement RAG, agent graph, tool calling, memory, API, backend, UI, logging, evaluation tests
Phase 5	Evaluate	Measure answer quality, RAG faithfulness, tool call accuracy, cost per request, latency, safety
Phase 6	Deploy	Cloud hosting, API gateway, containerisation, CI/CD, monitoring, user access, rollout plan
Phase 7	Monitor	Track token usage, cost per user, latency, error rate, hallucination, cache hits, user feedback
Phase 8	Improve	Prompt optimisation, model routing, semantic caching, better chunking, feedback loop, fine-tuning

Interfaces (Types you can Deliver)

Interface Type	Best for	Example
Chat Interface	Q&A, support, ^{internal} knowledge	"Ask our policy bot"
Form-based AI Tool	Structured ip → AI op	proposal generator
Dashboard + AI Insights	Data heavy clients	Sales forecasting UI
Voice Agents	Call centres, field workers	AI phone receptionist
✓ Embedded widget	Existing web/site/app integration	AI help button
Autonomous Agent portal	power users who want AI to Act	Research + Report Agent

8. Client Expectation Management

Always align expectations BEFORE writing a single line of code. This protects both you and the client.

Accuracy Expectations by Use Case

Use Case	MVP Accuracy Range	Risk Level
Internal FAQ / RAG Assistant	75-90%	Low -- great starting point
Structured Document Extraction	80-95%	Low -- good for invoices and forms
Support Draft Generation	70-90%	Medium -- human review recommended
Autonomous Tool Execution	60-85%	Medium -- start with approval gates
Complex Multi-Agent Reasoning	60-80%	Medium-High -- iterate aggressively
Legal / Financial Decisions	50-70%	High -- always human-in-the-loop

Human-in-the-Loop Rules (Default)

- Sending emails or messages on behalf of a user (webhook, SNS(AWS))
- Modifying records, databases, or production systems
- Financial transactions or approvals
- Legal, compliance, or medical output that will be acted upon
- Deleting or overwriting data
- Any customer-facing communication in the first four weeks

Success Metrics to Define Upfront

Metric	Example Target
Support ticket deflection	>= 30% reduction in human-handled tickets
Document search time	>= 50% reduction vs manual search
Cost per LLM request	< \$0.05 per query after optimisation
RAG faithfulness score	>= 0.80 (measured via RAGAS or DeepEval)
p95 response latency	< 10 seconds for 95% of requests
Human escalation rate	< 30% after 4 weeks of tuning
Prompt cost reduction	30 to 50% after prompt optimisation
Debug time reduction	>= 80% with tracing enabled

Recommendation Tech Stack for Consultants

Layer	Tool options
Agent building	Abacus AI, Langchain/graph, Crew AI, Auto Gen, Semantic Kernel, Bedrock Agent, Azure AI Agent Service
UI/Interface	Streamlit, Bubble, webflow + API
Workflow Automation	n8n, make.com, Zapier
Knowledge Base	Notion, Confluence + RAG pipeline
Hosting & Deployment	AWS, Vercel, Abacus AI
Client portal	custom web App (or) white-label SaaS

9. Solution Architectures

Architecture 1 -- Internal RAG Assistant

Best for: Internal knowledge bot, HR policy assistant, SOP search, document Q&A, training assistant.

Flow

Documents / PDFs -> Loader -> Chunking -> Embeddings -> Vector DB -> Retriever -> LLM -> Answer + Sources -> Observability

Component	Low Cost Option	Enterprise Option
LLM	GPT-4o mini / Gemini Flash / Claude 3 Haiku	GPT-4o / Claude 3.5 Sonnet / Gemini 1.5 Pro
Embeddings	text-embedding-3-small	text-embedding-3-large
Vector DB	Chroma / pgvector / Qdrant	Pinecone / OpenSearch / Azure AI Search
Framework	LangChain	LangGraph for more control
Observability	LangFuse (open source)	LangSmith for LangChain teams
Hosting	Cloud Run / Lambda / Azure Functions	Kubernetes / ECS / AKS for scale

Scale	Est. Monthly Cost
Prototype	\$20 - \$200
Small Business	\$200 - \$1,000
Mid-Scale	\$1,000 - \$5,000
Enterprise	\$5,000 - \$20,000+

Architecture 2 -- Tool-Calling Business Agent

Best for: CRM assistant, support triage, invoice lookup, order status, database query, operations assistant.

Flow

User Request -> Intent Detection -> LLM with Tools -> Tool Registry -> APIs / DB / SaaS -> Tool Result -> LLM Final Response -> Trace + Cost + Tool Logs

Component	Options
Agent Framework	LangGraph / Semantic Kernel / Bedrock Agents / Azure AI Agent Service
Tools	Lambda / Cloud Functions / Azure Functions / FastAPI endpoints
Memory / State	Redis / DynamoDB / Firestore / Cosmos DB
API Layer	FastAPI / API Gateway / APIM
Observability	LangFuse / LangSmith
Guardrails	Bedrock Guardrails / Azure Content Safety / Custom PII filters

Scale	Est. Monthly Cost
Prototype	\$100 - \$500
Small Business	\$500 - \$2,500
Mid-Scale	\$2,500 - \$10,000
Enterprise	\$10,000 - \$50,000+

Architecture 3 -- Supervisor Multi-Agent System

Best for: Research workflows, proposal generation, compliance review, complex support triage, content pipelines.

Flow

User Goal -> Supervisor Agent -> Task Planning -> Specialist Agents (Research / Data / Writer / Reviewer /

Tool] -> Synthesis -> Human Approval -> Final Output -> Observability + Feedback

- Use Supervisor Pattern -- specialists do not call each other directly
- Every handoff must go through the supervisor agent
- Use reducers to merge agent state at each step
- Use checkpointing for pause, resume, and human approval gates
- Wrap all agent outputs with deterministic validators
- Log every tool call, agent step, and model used

Scale	Est. Monthly Cost
Prototype	\$500 - \$2,000
Small Production	\$2,000 - \$8,000
Mid-Enterprise	\$8,000 - \$30,000
Large Enterprise	\$30,000 - \$150,000+

Architecture 4 -- Enterprise Agentic AI Platform

Best for: Large organisations wanting multiple agents across departments with full governance and observability.

Component	AWS	GCP	Azure
LLM	Bedrock (Claude/Titan)	Vertex AI (Gemini)	Azure OpenAI (GPT-4o)
Agent Framework	Bedrock Agents	Vertex ADK / Agent Builder	Semantic Kernel / AI Agent Svc
Tools	Lambda	Cloud Functions	Azure Functions
Workflow	Step Functions	Cloud Workflows	Logic Apps
Vector DB	OpenSearch / Bedrock KB	Vertex AI Vector Search	Azure AI Search
Memory	DynamoDB + ElastiCache	Firestore + Memorystore	Cosmos DB + Redis
Monitoring	CloudWatch + X-Ray	Cloud Monitoring + Trace	Azure Monitor + App Insights
Guardrails	Bedrock Guardrails	Vertex AI Safety	Azure Content Safety
Cost Control	AWS Cost Explorer	Cloud Billing	Azure Cost Management

10. LLM Model Cost Comparison (2025)

Output tokens cost 2 to 5x more than input tokens. Concise responses save money. Always route cheap tasks to cheap models.

Model	Provider	Input per 1M	Output per 1M	Best For
GPT-4o	OpenAI / Azure	\$2.50	\$10.00	General reasoning, best quality
GPT-4o mini	OpenAI / Azure	\$0.15	\$0.60	High-volume, cost-sensitive tasks
o1	OpenAI / Azure	\$15.00	\$60.00	Advanced reasoning problems
o3-mini	OpenAI / Azure	\$1.10	\$4.40	Efficient mid-tier reasoning
Claude 3.5 Sonnet	Anthropic / AWS	\$3.00	\$15.00	Best quality/cost balance
Claude 3.5 Haiku	Anthropic / AWS	\$0.80	\$4.00	Fast, affordable agent tasks
Claude 3 Haiku	Anthropic / AWS	\$0.25	\$1.25	Cheapest Claude option
Gemini 1.5 Pro	Google / GCP	\$1.25	\$5.00	Long-context tasks (1M tokens)
Gemini 1.5 Flash	Google / GCP	\$0.075	\$0.30	Very fast, very cheap
Gemini 2.0 Flash	Google / GCP	\$0.10	\$0.40	Latest fast model
Mistral Large	Mistral / AWS	\$2.00	\$6.00	EU data residency
Llama 3.1 70B	Self-hosted	~\$0.50-\$1 infra	N/A	Open-source, private deploy

Model Routing Strategy

Task Type	Recommended Model
Classification / Extraction	GPT-4o mini / Gemini Flash / Claude 3 Haiku
Simple Q&A / Summarisation	GPT-4o mini / Gemini Flash / Claude Haiku
Complex Reasoning	GPT-4o / Claude 3.5 Sonnet / Gemini 1.5 Pro
Code Generation	GPT-4o / Claude 3.5 Sonnet / o3-mini
Long Document Analysis	Gemini 1.5 Pro or Flash (1M context)
Legal / Compliance Review	GPT-4o / Claude 3.5 Sonnet + human approval
Repeated / Cached Queries	Semantic cache -- no LLM call needed

11. Cloud Infrastructure and Cost Structure

Compute Pricing by Instance Type (2025)

Instance	AWS per hr	GCP per hr	Azure per hr
4 vCPU / 16 GB CPU	\$0.23 (m5.xl)	\$0.19 (n1-std-4)	\$0.27 (DS3v2)
8 vCPU / 32 GB CPU	\$0.46 (m5.2xl)	\$0.38 (n1-std-8)	\$0.54 (DS4v2)
T4 GPU 16 GB	\$0.74 (g4dn.xl)	\$0.35 (n1+T4)	\$0.90 (NC4as)
V100 GPU 16 GB	\$3.83 (p3.2xl)	\$2.48 (n1+V100)	\$3.06 (NC6s)
A100 GPU 40 GB	\$32.77 (p4d.xl)	\$3.67 (a2-high)	\$3.40 (NC24ads)
H100 GPU 80 GB	~\$98/hr	~\$12/hr	~\$15/hr

Key Managed Services Pricing

Service	Approx. Pricing
AWS Bedrock Claude 3 Haiku	\$0.25/1M input + \$1.25/1M output
AWS Bedrock Claude 3.5 Sonnet	\$3.00/1M input + \$15.00/1M output
AWS OpenSearch Serverless	\$0.24/OCU-hr -- vector DB for RAG
AWS Step Functions	\$0.025/1K state transitions
AWS Lambda	\$0.20/1M requests + compute
GCP Vertex Gemini 1.5 Flash	\$0.075/1M input + \$0.30/1M output
GCP Vertex AI Vector Search	\$0.09/hr per node
GCP Pub/Sub	\$0.04/GB messaging
Azure OpenAI GPT-4o	\$2.50/1M input + \$10.00/1M output
Azure AI Search S1	\$0.10/hr -- vector + keyword search
Azure Logic Apps	\$0.000025 per action -- workflow automation

Monthly Cost Estimates by Solution

Solution	Total per Month	LLM Tokens	Compute	Observability
Small RAG Assistant	\$90 - \$1,100	\$50 - \$500	\$20 - \$200	\$0 - \$100
Tool-Calling Agent	\$420 - \$3,700	\$200 - \$2,000	\$100 - \$500	\$50 - \$300
Multi-Agent Workflow	\$2,000 - \$25,000	\$1,000 - \$15,000	\$500 - \$2,000	\$200 - \$2,000
Enterprise AI Platform	\$10K - \$180K+	\$5K - \$100K+	\$2K - \$20K	\$500 - \$10K

12. Observability and Cost Management ROI

Core Equation

Observability = Money Saved

50 to 80 percent cost reduction is achievable. The question is not "can we afford observability?" It is "can we afford NOT to have it?"

ROI Model Example

Cost Category (Without Observability)	Amount
Monthly LLM Spend	\$20,000
Estimated Token Waste (30%)	\$6,000
Debug Time (10 hrs/week x \$100)	\$4,000/month
Incident Cost (amortised)	\$5,000
TOTAL PREVENTABLE COST	\$15,000/month

With Observability	Value
Platform Cost (LangFuse/LangSmith)	\$500 - \$2,000/month
Setup Time (one-time only)	~8 hours
ROI in First Month	7x to 30x return
Token Cost Reduction	~47% via prompt optimisation
Debug Time Reduction	~80% with tracing enabled
Cost Surprises Prevention	\$0 runaway surprises with alerts

Savings Formula

Savings = Token Waste + (Debug Time x Hourly Rate) + (Incidents Prevented x Avg Incident Cost)

Example: \$6,000 + \$4,000 + \$5,000 = \$15,000/month saved for only \$500 - \$2,000 investment.

The 3 Cost Levers

Lever	Savings	Effort	How
1. Prompt Optimisation	30-50%	Low	Reduce tokens, concise instructions, fewer examples in prompt
2. Semantic Caching	30-50%	Medium	Cache similar queries -- no LLM call for repeated questions
3. Model Routing	50-70%	Medium	Route simple tasks to cheap models, complex to premium
All Three Combined	70-85%	Medium	Maximum ROI: router + cache + optimised prompts together

13. Proposal and Delivery Structure

Proposal Document Sections

Section	Contents
1. Executive Summary	Problem, proposed solution, expected business impact
2. Current Pain	Detailed description of the workflow inefficiency
3. Proposed Solution	Agent type, capabilities, boundaries, human approval rules
4. Architecture	LLM, RAG, tools, cloud, DB, vector DB, observability, security
5. Scope	What is included -- and critically, what is excluded
6. Timeline	Week-by-week milestones with clear deliverables per milestone
7. Success Metrics	Specific, measurable targets agreed before work begins
8. Cost Estimate	Service fee + cloud cost + LLM cost + maintenance cost
9. Risks and Assumptions	Data access, API availability, compliance, accuracy expectations
10. Commercial Terms	Payment schedule, change request policy, support SLA

Sample 6-Week MVP Timeline

Week	Focus	Key Deliverables
Week 1	Discovery + Architecture	Finalised requirements, architecture diagram, data access
Week 2	Data Ingestion + RAG	Document loading, chunking, embedding, vector DB setup
Week 3	Agent + Tool Calling	LLM agent graph, tool functions, memory, session handling
Week 4	UI/API + Observability	User interface or API, LangFuse tracing, cost dashboard
Week 5	Testing + Evaluation	RAGAS evaluation, edge case testing, accuracy report
Week 6	Deployment + Handover	Cloud deploy, documentation, client training, final delivery

Final Delivery Package

- Product requirements document
- Architecture diagram and technical documentation
- Source code repository with README
- Deployment and environment setup guide
- API documentation
- User guide and admin guide
- Evaluation report: accuracy, latency, and cost
- Observability dashboard (LangFuse or LangSmith)
- Cost breakdown and ROI report
- Security and compliance notes
- Improvement roadmap for next 3 months

Payment Structure

Milestone (Small MVP)	Payment	Trigger
Project Start	50%	Signed scope, kick off
MVP Demo	30%	Working prototype demonstrated
Final Handover	20%	All docs and deployment complete

Milestone (Large Project)	Payment	Trigger
Discovery + Architecture	20%	Scope signed, architecture approved
Prototype	25%	Core agent demonstrated
Integration	25%	Connected to real business systems
Testing + Deployment	20%	Evaluation report and deployed system
Handover	10%	Full documentation delivered

14. Monetisation Model and Retainer Strategy

Move every client from a one-time project to a monthly retainer. Retainers are your recurring revenue engine.

Level	Offer	Price	Purpose
Level 1	AI Workflow Audit	\$500 - \$2,000	Entry offer: get paid to discover
Level 2	MVP Build	\$3,000 - \$15,000	Prove value and build trust
Level 3	Production Build	\$15,000 - \$75,000	Deploy business-critical system
Level 4	Enterprise Multi-Agent Platform	\$50,000 - \$150,000+	Multi-agent, full cloud, governance
Level 5	Monthly Retainer	\$1,000 - \$10,000/mo	Maintenance, monitoring, optimisation

What Retainer Services Include

- Monthly LLM cost optimisation report (prompts, routing, caching)
- Prompt engineering updates as model behaviour evolves
- Observability dashboard review: latency, errors, costs
- Incident debugging and resolution
- New tool or integration development
- Security and compliance updates
- Evaluation dataset maintenance and quality scoring
- User feedback review and agent improvement
- Quarterly architecture review and roadmap update

Retainer Pricing by Client Size

Client Size	Monthly Price	Effort	Scope
Startup / Small Business	\$1,000 - \$2,500/month	5-10 hrs/month	Basic monitoring + optimisation
Mid-Market Company	\$2,500 - \$5,000/month	10-20 hrs/month	Full optimisation + improvement sprints
Enterprise	\$5,000 - \$10,000+/month	20-40 hrs/month	Ongoing development + governance

15. Your 90-Day Business Launch Plan

Period	Focus	Key Actions
Days 1-15	Foundation	Pick 2-3 target industries. Create service offer + pricing. Build proposal + ROI calculator templates. Build one working demo agent. Set up GitHub portfolio and LinkedIn positioning.
Days 16-30	Demo Portfolio	Build Demo 1: RAG assistant over company documents. Build Demo 2: Tool-calling agent with fake CRM. Each demo shows: user query, sources, tool calls, cost per request, trace dashboard, error handling.
Days 31-60	Outreach + Pitching	Contact 100 companies, founders, or department managers. Offer 20-minute AI Workflow Audit. Run discovery calls using the discovery framework. Send 5-10 tailored proposals. Close first paying MVP.
Days 61-90	Deliver + Expand	Deliver first MVP with full documentation. Measure and report impact vs agreed metrics. Create case study with numbers and testimonial. Pitch observability retainer to same client. Start outreach to 3 similar companies.

Weekly Time Allocation

Activity	Allocation	40-Hour Week
Lead Generation and Outreach	25%	~10 hrs
Discovery Calls and Proposals	15%	~6 hrs
Building Client Solutions	35%	~14 hrs
Learning and Research	10%	~4 hrs
Content and Authority Building	10%	~4 hrs
Admin and Finance	5%	~2 hrs

16. Final Mental Model and Differentiation

Enterprise AI Stack -- Your Knowledge Map

Layer	Purpose	Key Technologies
Data Layer	Collect, store, and govern all enterprise data	Kafka, S3/GCS/ADLS, Snowflake, Delta Lake
Feature Layer	Transform raw data into ML-ready features	Feast, Tecton, Spark, dbt
Model Layer	Train, evaluate, version, and serve ML models	PyTorch, XGBoost, MLflow, SageMaker
LLM Layer	Reasoning, language understanding, generation	GPT-4o, Claude, Gemini, Llama
Agent Layer	Orchestrate tools, models, and workflows autonomously	LangGraph, CrewAI, Bedrock Agents, ADK
Execution Layer	Run deterministic actions: APIs, DBs, workflows	Lambda, Step Functions, Cloud Run
Observability Layer	Monitor everything: infra, models, agents, costs	Grafana, LangSmith, LangFuse, Evidently AI
Governance Layer	Security, compliance, audit, access control	IAM, KMS, Purview, Bedrock Guardrails
Feedback Layer	Close the loop: retrain, fine-tune, and improve	RLHF, Label Studio, MLflow, Airflow

Your Differentiation Pillars

Pillar	What It Means for the Client
Business-First Discovery	You start with workflow ROI, not technology selection
Product Manager Thinking	You define scope, metrics, risks, and non-goals before building
Production-Grade Architecture	You design for reliability, security, and observability from day one
Multi-Agent Expertise	You understand supervisor patterns, state management, and checkpointing
Observability and Cost Control	You can prove and reduce LLM costs -- most consultants cannot
Cloud Infrastructure Awareness	You speak AWS, GCP, and Azure fluently
Clear Handover and Docs	Clients can maintain and extend your work independently
Retainer-Based Optimisation	You provide ongoing value, not just one-time delivery

Client Conversation Scripts

Client: "We want AI."

You: Before choosing a model, I would like to understand which workflow is slow, expensive, or error-prone. The best AI projects start from workflow ROI, not technology selection.

Client: "Can you build a chatbot?"

You: Yes -- but first let us clarify: do you need simple Q&A, a RAG assistant, a tool-calling agent, or a multi-agent workflow? The right architecture depends on whether it only answers questions or actually performs business tasks.

Client: "Will it be accurate?"

You: We improve accuracy using RAG, citations, guardrails, evaluation datasets, and human approval gates. We should define an acceptable accuracy target and escalation rules -- not assume 100% automation from day one.

Client: "How much will it cost?"

You: There are two costs: build cost and running cost. Running cost depends on LLM tokens, users, vector DB, and cloud hosting. I estimate both during the architecture phase and optimise with caching, prompt reduction, and model routing.

The Final Formula

Business Pain

- + AI Product Thinking
- + Agentic Architecture
- + Cloud Infrastructure
- + Observability and Cost Control
- + Clear Delivery and Documentation
- = A Valuable, Sustainable Agentic AI Consulting Business

"This person does not just know AI tools. They understand business, product, infrastructure, cost, risk, and delivery."