

AI / ML CONSULTING FRAMEWORK

Complete Consultant's Playbook

Approach · Pitch · Deliver · Monetise · Scale

Traditional ML · LLMs · MLOps · Cloud Architecture · Cost Structure

IDENTITY

AI/ML Consultant + Product Builder

EARN

\$5K to \$200K+ per engagement

BUILD

ML Pipelines, LLM Products, MLOps

STACK

AWS, GCP, Azure + Open Source

01 · CONSULTING IDENTITY & POSITIONING

Who You Are · What You Sell · Why Clients Pay Premium

CORE IDENTITY Your Role as an AI/ML Consultant

You translate business problems into production-grade machine learning systems — from raw data pipelines to deployed models, LLM products, MLOps infrastructure, and measurable business outcomes. You are part scientist, part engineer, part product manager.

Three Hats You Wear

Hat	Role	What You Deliver
ML Scientist	Model Builder	Trained, evaluated, production-ready ML models
ML Engineer	System Architect	End-to-end ML pipelines, feature stores, serving infra
Product Specialist	AI Product Builder	LLM-powered products, APIs, dashboards, ROI stories

Your Value Proposition

- Speed: Deliver in weeks what in-house teams take months to build
- Breadth: Data to Model to Deploy to Monitor — full lifecycle ownership
- ROI Focus: Every engagement tied to cost savings, revenue lift, or efficiency gains
- Cloud Agnostic: AWS Bedrock, GCP Vertex, Azure OpenAI, Open-Source stack
- LLM + Classical ML: Rare combo — most consultants specialise in only one

02 · THE AI/ML CONSULTING FLYWHEEL

10-Step Repeatable Process from Cold Lead to Retained Client

Step	Stage	Action
1	Identify	Find company with an AI/ML pain point or unused data asset
2	Outreach	Personalised cold email/LinkedIn with a specific AI use-case hypothesis
3	Discovery	60-min discovery call: understand data, goals, blockers, budget
4	Audit	Free 2-hr data/model audit — find quick wins and scope engagement
5	Proposal	Scoped proposal: problem, solution, timeline, price, ROI
6	Kickoff	Data access, environment setup, weekly cadence established
7	Build	Sprint-based delivery: model, evaluation, API, dashboard
8	Present	Present results with business metrics, not just model metrics
9	Handover	Documentation, training, runbooks, knowledge transfer
10	Retainer	Monthly retainer for monitoring, retraining, enhancements

FLYWHEEL SECRET How the Business Compounds

Each delivered project becomes a case study. Case studies attract better clients. Better clients pay higher retainers. Retainers fund your time for thought leadership. Thought leadership fills your pipeline. That is the flywheel.

03 · IDEAL CLIENT PROFILE & QUALIFYING SIGNALS

Company Type	Size	AI Readiness	Sweet Spot Problem
Series B-D Startups	50-500 employees	Has data, no ML team	First production ML model
Mid-Market SaaS	100-1000 employees	Has analysts, no MLEs	Churn prediction / LLM feature
Enterprise Division	1000+ employees	Has data science, no MLOps	MLOps infra and cost reduction
Traditional Industry	Any	Digitising fast	Process automation with ML/AI
VC-backed AI Startup	10-100 employees	Building AI core	LLM product architecture

Green Flags — Prioritise These Leads

- Job postings for ML Engineer, Data Scientist — they want AI but can't hire fast enough
- C-suite posting on LinkedIn about AI strategy or digital transformation
- Company blog posts about data challenges or failed AI pilots
- Using Snowflake, dbt, Databricks — data-mature and ML-ready
- Recent funding round — budget available for AI investment

Red Flags — Qualify Out Early

- No structured data and no data team — too early, too long of an engagement
- Just want ChatGPT on their data with no clear use case
- No internal champion to push decisions through stakeholders

04 · OUTREACH, PITCH & DISCOVERY FRAMEWORK

Cold Outreach Template

Subject: Quick idea for [Company] + predictive ML

Hi [Name],

I noticed [Company] is [doing X / hiring Y / posting about Z]. I help [industry] companies build production ML systems such as churn prediction, demand forecasting, and LLM-powered features in 6-10 weeks.

One pattern I often see with companies at your stage: [specific pain point hypothesis]. Happy to share a quick framework that has worked for similar companies with no pitch, just value.

Would a 20-minute call this week make sense?

Best, [Your Name]

Discovery Call Question Framework

Phase	Question	What You Are Learning
Business	What is the #1 metric your team is measured on?	True north star KPI
Business	What would a 10% improvement in that metric mean?	ROI potential in dollars
Data	What data do you currently collect and where does it live?	Data readiness assessment
Data	How long is your historical data? Any labelled outcomes?	ML feasibility check
Technical	Do you have a data/ML team? What is their stack?	Build vs. guide decision
Technical	What have you tried before? What did not work?	Avoid repeating failures
Process	What does a solved version of this problem look like?	Success criteria definition
Budget	Is there a budget allocated for AI/ML this year?	Deal qualification signal
Timeline	Is there a board or exec deadline driving this?	Urgency and priority level
Politics	Who else needs to sign off on this decision?	Stakeholder map

05 · AI/ML OPPORTUNITY SCORING MATRIX

Score each use case 1 to 5 on each dimension. Total score 20 or above = pursue. 28 or above = top priority.

Dimension	Score 1 (Low)	Score 3 (Medium)	Score 5 (High)
Data Availability	No structured data	Some data with gaps	Clean labelled data ready
Business Impact	Under \$50K potential	\$50K to \$500K potential	Over \$500K or strategic
Technical Feasibility	Unproven ML approach	Known ML pattern	Solved problem, clear path
Time to Value	Over 6 months	3 to 6 months	Under 3 months
Stakeholder Buy-in	No champion identified	1 champion exists	C-suite backed
Competitive Moat	Anyone can copy	Some data advantage	Data moat / network effect

TOP ML USE CASES BY INDUSTRY Quick Reference for Qualifying Conversations

Retail: Demand forecasting, churn prediction, recommendation engine
Finance: Credit scoring, fraud detection, algorithmic trading signals
Healthcare: Patient readmission risk, diagnostic assist, drug discovery
SaaS: Churn and expansion prediction, feature adoption ML, support ticket routing
Manufacturing: Predictive maintenance, defect detection via computer vision, yield optimisation
Logistics: Route optimisation, ETA prediction, demand sensing

06 · SERVICE OFFERINGS & PRICING LADDER

Tier	Service	Timeline	Price Range	Deliverable
Starter	AI/ML Readiness Audit	1 week	\$2,000 - \$5,000	Report + roadmap + quick wins
Core	Single ML Model Build	4-6 weeks	\$8,000 - \$20,000	Trained model + API + docs
Advanced	End-to-End ML Pipeline	6-10 weeks	\$20,000 - \$50,000	Full pipeline + MLOps infra
Platform	AI Product Build (LLM)	8-16 weeks	\$30,000 - \$80,000	Deployed AI product + CI/CD
Enterprise	Full ML Platform + Team	3-6 months	\$80,000 - \$200,000+	Enterprise AI stack
Retainer	Monitoring and Improvement	Ongoing	\$3,000 - \$15,000/mo	Monthly reports + retraining

Pricing Strategy Tips

- Start with the Audit at \$2K-\$5K — low risk for client, high intelligence for you
- Anchor on ROI: if your model saves \$500K/year, a \$40K project = 12x ROI
- Value-based pricing beats hourly: charge for outcomes, not time
- Package retainers into every project delivery as the natural next step
- Offer 3-tier proposals: Basic, Standard, Premium — 80% pick the middle

07 · PROJECT LIFECYCLE — 8-PHASE DELIVERY FRAMEWORK

Phase	Name	Duration	Key Activities	Outputs
1	Discovery and Data Audit	Week 1	EDA, data quality, feasibility check	Audit report, go/no-go decision
2	Problem Framing	Week 1-2	Define ML task, metrics, baseline	ML specification document
3	Data Pipeline Build	Week 2-3	Ingestion, transformation, feature engineering	Production data pipeline
4	Model Development	Week 3-5	Train, evaluate, tune, compare models	Best model + evaluation report
5	Model Serving	Week 5-6	API wrapping, containerisation, deploy	Live model endpoint
6	Monitoring Setup	Week 6-7	Dashboards, alerts, drift detection	MLOps monitoring stack
7	Handover and Training	Week 7-8	Docs, runbooks, team training	Knowledge transfer package
8	Retainer Kickoff	Week 8+	Monthly cadence, retraining plan	Ongoing ML operations

08 · SOLUTION ARCHITECTURES & COST STRUCTURE

Architecture 1: Classic ML Model (Batch Prediction)

Layer	Technology	AWS Option	GCP Option	Azure Option	Monthly Cost
Data Storage	S3 / GCS / ADLS	S3 Standard	GCS Standard	ADLS Gen2	\$20-100
Data Pipeline	Airflow / dbt	MWAA	Cloud Composer	ADF	\$200-800
Feature Store	Feast / Redis	SageMaker FS	Vertex FS	AzureML FS	\$100-400
Model Training	sklearn / XGBoost	SageMaker Training	Vertex Training	AzureML Compute	\$100-500
Batch Serving	Lambda / Cloud Run	Batch Transform	Batch Prediction	Batch Endpoints	\$50-300
Monitoring	Evidently + Grafana	Model Monitor	Vertex Monitor	AzureML Monitor	\$100-300
TOTAL					\$570 - \$2,400/mo

Architecture 2: Real-Time ML API (Online Prediction)

Layer	Technology	AWS	GCP	Azure	Monthly Cost
Data Ingestion	Kafka / Kinesis	Kinesis	Pub/Sub	Event Hubs	\$100-500
Feature Store Online	Redis + Feast	SageMaker FS Online	Vertex FS Online	AzureML FS	\$200-600
Model Training	PyTorch / sklearn	SageMaker	Vertex AI	AzureML	\$200-800
Real-Time Serving	FastAPI + Docker	SageMaker Endpoint	Vertex Endpoint	AzureML Endpoint	\$300-1,500
API Gateway	Kong / AWS APIGW	API Gateway	Apigee	Azure APIM	\$50-200
Monitoring	Prometheus + Grafana	CloudWatch	Cloud Monitoring	Azure Monitor	\$100-300
TOTAL					\$950 - \$3,900/mo

Architecture 3: LLM-Powered AI Product (RAG + Fine-tuning)

Layer	Technology	AWS	GCP	Azure	Monthly Cost
-------	------------	-----	-----	-------	--------------

LLM API	GPT-4o / Claude / Gemini	Bedrock	Vertex AI	Azure OpenAI	\$500-5,000
Vector DB	Pinecone / Weaviate	OpenSearch	Vertex Vector Search	Azure AI Search	\$100-500
Embedding Model	OpenAI / Cohere	Bedrock Titan	Vertex Embeddings	Azure OpenAI Ada	\$50-500
Orchestration	LangChain / LangGraph	Lambda	Cloud Run	Azure Functions	\$50-300
LLM Observability	LangSmith / Helicone	LangSmith	LangSmith	LangSmith	\$0-200
App Hosting	FastAPI + Docker + K8s	EKS / ECS	GKE	AKS	\$200-800
TOTAL					\$900 - \$7,300/mo

Architecture 4: Enterprise ML Platform (Full MLOps)

Layer	Technology	AWS	GCP	Azure	Monthly Cost
Data Lake	Delta Lake / Iceberg	S3 + Glue	GCS + BigLake	ADLS + Synapse	\$200-1,000
Data Pipelines	Spark + Airflow	EMR + MWAA	Dataproc + Composer	Synapse + ADF	\$500-2,000
Feature Store	Feast + Redis	SageMaker FS	Vertex FS	AzureML FS	\$200-800
Experiment Tracking	MLflow / W&B	MLflow on SM	Vertex Experiments	AzureML Experiments	\$0-500
Training GPU	PyTorch + DDP	SageMaker p3/p4	Vertex A100	AzureML NC-series	\$500-3,000
Model Serving	Triton + BentoML	Multi-Model Endpoints	Vertex Endpoints	AzureML Endpoints	\$300-1,500
CI/CD	GitHub Actions + ArgoCD	CodePipeline	Cloud Build	Azure DevOps	\$100-400
Drift Detection	Evidently AI	Model Monitor	Vertex Monitor	AzureML Monitor	\$100-400
Governance	IAM + KMS + VPC	GuardDuty + Macie	Security Command	Defender + Purview	\$200-600
TOTAL					\$2,100 - \$10,200/mo

09 · LLM & ML MODEL COST REFERENCE (2025)

Model	Provider	Input \$/1M	Output \$/1M	Best Use Case
GPT-4o	OpenAI / Azure	\$2.50	\$10.00	General reasoning, complex tasks
GPT-4o mini	OpenAI / Azure	\$0.15	\$0.60	High-volume, cost-sensitive
o1	OpenAI / Azure	\$15.00	\$60.00	Complex multi-step reasoning
o3-mini	OpenAI / Azure	\$1.10	\$4.40	Efficient coding and reasoning
Claude 3.5 Sonnet	Anthropic / AWS	\$3.00	\$15.00	Best quality/cost balance
Claude 3.5 Haiku	Anthropic / AWS	\$0.80	\$4.00	Fast, affordable production
Claude 3 Haiku	Anthropic / AWS	\$0.25	\$1.25	Cheapest Claude option
Gemini 1.5 Pro	Google / GCP	\$1.25	\$5.00	Long-context (1M tokens)

Gemini 1.5 Flash	Google / GCP	\$0.075	\$0.30	Fast, long context, cheap
Gemini 2.0 Flash	Google / GCP	\$0.10	\$0.40	Latest fast model from Google
Llama 3.1 70B	Self-hosted	~\$0.50-1.00 infra	N/A	Open source, no API cost
Mistral Large	Mistral / AWS	\$2.00	\$6.00	European data residency

3-Tier Model Routing Strategy

Query Type	Complexity	Recommended Model	Cost Saving vs GPT-4o
Simple lookup / classification	Low	GPT-4o mini / Haiku / Flash	70-80% savings
Standard reasoning / generation	Medium	Claude 3.5 Haiku / GPT-4o mini	50-60% savings
Complex analysis / long-context	High	GPT-4o / Claude 3.5 Sonnet	Baseline cost
Multi-step math / deep reasoning	Very High	o1 / o3-mini	Premium - use sparingly
Batch / offline processing	Any	Self-hosted Llama / Mistral	90%+ savings

10 · CLOUD COMPUTE PRICING REFERENCE (2025)

Instance	vCPU	RAM	GPU	AWS \$/hr	GCP \$/hr	Azure \$/hr	Use Case
4 vCPU / 16 GB	4	16 GB	None	\$0.23	\$0.19	\$0.27	API serving, batch jobs
8 vCPU / 32 GB	8	32 GB	None	\$0.46	\$0.38	\$0.54	Feature engineering
T4 GPU 16 GB	4	16 GB	T4	\$0.74	\$0.35	\$0.90	Inference, fine-tuning
V100 GPU 16 GB	8	61 GB	V100	\$3.83	\$2.48	\$3.06	Model training
A100 GPU 40 GB	12	85 GB	A100	\$32.77	\$3.67	\$3.40	LLM fine-tuning
H100 GPU 80 GB	96	384 GB	H100	~\$98/hr	~\$12/hr	~\$15/hr	Foundation model training

Monthly Cost Estimates by Solution Scale

Scale	Description	AWS/mo	GCP/mo	Azure/mo	Key Cost Driver
POC / Prototype	1 model, low traffic, no HA	\$200-800	\$180-700	\$220-850	Compute only
Small Production	1-3 models, moderate traffic	\$800-3,000	\$700-2,700	\$850-3,200	Infra + LLM API
Mid-Scale	5-10 models, 1M preds/day	\$2,000-8,000	\$1,800-7,200	\$2,200-8,500	All layers active
Enterprise	Full platform, 10M+ preds/day	\$8,000-30,000+	\$7,000-27,000+	\$8,500-32,000+	GPU + storage + ops

COST OPTIMISATION QUICK WINS 6 Ways to Cut Cloud ML Costs Immediately

1. Spot / Preemptible instances for training jobs: up to 90% savings

2. Serverless inference for low-traffic models: zero idle compute cost
3. Semantic caching of LLM responses: 30-50% cost reduction
4. Quantised models (INT8/INT4) for inference: 2-4x cheaper GPU cost
5. Committed use / savings plans for baseline workloads: 30-40% discount
6. Model routing to cheap models for simple queries: 60-80% LLM savings

11 · MLOPS & OBSERVABILITY TOOLING REFERENCE

Category	Tool	Free Tier	Paid Tier	Best For
Experiment Tracking	MLflow (OSS)	Free	Infra cost only	Most popular open source
Experiment Tracking	Weights and Biases	Free personal	\$50+/user/mo	Best UI, team collab
Experiment Tracking	Comet ML	Free personal	\$179+/mo	Enterprise features
Feature Store	Feast (OSS)	Free	Infra cost only	Most popular OSS FS
Feature Store	Tecton	No free tier	Custom pricing	Enterprise feature store
LLM Observability	LangSmith	Free 100 traces/mo	\$39+/mo	Best for LangChain apps
LLM Observability	Arize Phoenix	Free OSS	\$500+/mo	Multi-framework support
LLM Observability	Helicone	Free 10K req/mo	\$20+/mo	Lightweight proxy
Drift Detection	Evidently AI	Free OSS	\$500+/mo cloud	Best OSS drift tool
Drift Detection	NannyML	Free OSS	Custom	No ground truth needed
LLM Evaluation	RAGAS	Free OSS	Infra cost	RAG pipeline evaluation
LLM Evaluation	DeepEval	Free OSS	\$49+/mo	Comprehensive LLM evals
Data Annotation	Label Studio	Free OSS	\$500+/mo	Multi-modal labelling
Model Serving	BentoML	Free OSS	Infra cost	Flexible model serving



12 · ENTERPRISE AI STACK — ULTIMATE MENTAL MODEL

Layer	What It Does	Key Technologies	Consultant Role
Data Layer	Collect, store, govern all data	Kafka, S3/GCS/ADLS, Snowflake, Delta Lake	Data Architect
Feature Layer	Transform raw data to ML features	Feast, Tecton, Spark, dbt	Feature Engineer
Model Layer	Train, evaluate, version, serve ML	PyTorch, XGBoost, MLflow, SageMaker	ML Engineer
LLM Layer	Reasoning, language, generation	GPT-4o, Claude, Gemini, Llama	LLM Engineer
Agent Layer	Orchestrate tools autonomously	LangGraph, CrewAI, Bedrock Agents	Agentic AI Builder
Execution Layer	Run APIs, databases, workflows	Lambda, Step Functions, Cloud Run	Backend Engineer
Observability Layer	Monitor models, agents, costs	Grafana, LangSmith, Evidently AI	MLOps Engineer
Governance Layer	Security, compliance, audit	IAM, KMS, Purview, Bedrock	ML Governance

		Guardrails	
Feedback Layer	Retrain, fine-tune, improve	RLHF, Label Studio, MLflow, Airflow	ML Scientist



THE CONSULTANT'S STACK CHEAT CODE Your Greatest Differentiator

Know all 9 layers. Clients only know 1-2. Your ability to speak across all layers is your greatest differentiator. When you walk into a meeting and map their problem to the right layer instantly, they see you as the expert they have been looking for.

13 · CLIENT EXPECTATION MANAGEMENT

What Clients Expect	Reality	How to Frame It
100% accurate model	80-95% is excellent production quality	Baseline vs. model lift narrative
AI replaces all human decisions	AI augments, humans approve high-risk calls	Human-in-the-loop design
Instant results	6-10 weeks minimum for production	Phase-by-phase milestone plan
One-time project, done forever	Models drift, need retraining over time	Retainer as maintenance contract
AI works on any data	Garbage in, garbage out — data quality first	Data quality audit as Phase 1
Cheaper than hiring	More expensive short-term, cheaper long-term	3-year TCO, not 3-month cost

THE #1 CONSULTING MISTAKE Optimise for Business Outcome, Not Model Accuracy

A model that is 72% accurate but saves \$2M per year is infinitely better than a 95% accurate model that moves no KPI. Always anchor the engagement on business metrics, not ML metrics. Define success criteria in dollar terms before writing a single line of code.

14 · PROPOSAL & DELIVERY STRUCTURE

Section	Content	Length
1. Executive Summary	Problem, Solution, ROI in 3 bullet points	Half page
2. Problem Statement	Their words, their pain, their cost — reflect back discovery	1 page
3. Proposed Solution	Architecture diagram + phased plan + tech stack	2 pages
4. Timeline and Milestones	Week-by-week Gantt with clear deliverables	1 page
5. Investment and Terms	3-tier pricing + payment schedule + retainer option	1 page

Final Delivery Package Checklist

- Trained model artifact (pickled / ONNX / container image)
- Model card: performance metrics, training data summary, known limitations
- REST API endpoint with OpenAPI/Swagger documentation
- Data pipeline code in Git repo with full README
- MLflow or W&B experiment tracking workspace access
- Monitoring dashboard with alert configurations
- Runbook: how to retrain, redeploy, and rollback
- Business impact report: before vs. after metrics with dollar impact
- Retainer proposal for ongoing operations and improvements

Payment Structure

Milestone	When	Percentage
Kickoff Payment	On contract signing	30%
Mid-Project Milestone	On delivery of working model	40%
Final Delivery	On acceptance of full deliverable package	30%
Retainer	Monthly, billed in advance	100% monthly

15 · MONETISATION MODEL & REVENUE LADDER

Revenue Stream	Description	Price	Scalability
ML Audit	Entry-point engagement	\$2K - \$5K	Low (1:1)
Project-Based	Fixed-scope ML builds	\$10K - \$80K	Medium
Retainer	Monthly MLOps and improvements	\$3K - \$15K/mo	High recurring
LLM Product Build	AI product development	\$30K - \$100K	Medium
Training / Workshop	Team upskilling sessions	\$3K - \$10K/day	High group
Fractional Head of AI	Part-time strategic leadership	\$5K - \$20K/mo	High
White-label Platform	Sell your AI stack to agencies	Custom	Very High
AI SaaS Product	Productise your consulting IP	\$99 - \$999/mo	Highest passive

TARGET REVENUE MILESTONES Build Toward a Predictable Monthly Run Rate

Month 1-3: \$5K-\$15K/month (1 audit + 1 small project)

Month 4-6: \$15K-\$30K/month (1 project + 1 retainer)

Month 7-12: \$30K-\$60K/month (2 projects + 2-3 retainers)

Year 2+: \$60K-\$150K/month (retainers + training + productised services)

16 · 90-DAY AI/ML CONSULTING BUSINESS LAUNCH PLAN

Phase	Days	Focus	Key Actions	Success Metric
Foundation	1-15	Identity and	Define niche, update	Clear ICP defined,

		Positioning	LinkedIn, build 1-page website, create case study template	digital presence live
Proof Work	16-30	Build Credibility	Do 1 free or discounted audit, document results, write LinkedIn case study post	1 published case study
Pipeline	31-45	Outreach Engine	50 personalised LinkedIn DMs per week, join 3 AI communities, attend 1 meetup	10 discovery calls booked
First Revenue	46-60	Close First Deal	Convert 1 audit to project, negotiate scope, get deposit paid	\$5K-\$20K first invoice paid
Systematise	61-75	Process and Templates	Build proposal template, onboarding checklist, delivery framework docs	Repeatable delivery system built
Scale	76-90	Pipeline and Retainer	Close 2nd project, upsell retainer to first client, consider first subcontractor	\$15K+ per month run rate

17 · DIFFERENTIATION PILLARS & CLIENT SCRIPTS

Pillar	What It Means	How to Communicate It
Full-Stack ML	Data to Model to Deploy to Monitor	I own the entire ML lifecycle, not just the model
Business-First	ROI framing on every deliverable	Every model I build has a dollar value attached
Cloud Agnostic	AWS + GCP + Azure + Open Source	I fit your existing stack, not the other way around
LLM + Classical ML	Both modern and proven approaches	I know when to use LLMs and when not to
Speed to Value	Production in weeks, not months	POC in week 2, production by week 8

WHEN ASKED: Why Not Just Hire a Data Scientist? Script to Overcome the Build-vs-Buy Objection

A data scientist costs \$150K-\$200K per year in salary plus benefits and onboarding time. I will deliver a production-ready ML system in 8 weeks for a fraction of that cost. After delivery your team owns it completely — and I can stay on as a fractional advisor for \$3K/month.

WHEN ASKED: What Makes Your Models Better? Script to Elevate the Technical Conversation

Model accuracy is the last 20% of the problem. The first 80% is data quality, feature engineering, and deployment infrastructure. I specialise in the full 100%. That is why my models stay accurate in production, not just in notebooks.

WHEN NEGOTIATING PRICE ROI Reframe Script

Let us look at the ROI together. If this churn model retains 2% more customers at \$500 ARR each, and you have 10,000 customers — that is \$1M in saved ARR annually. My fee is \$25,000. That is a 40x return in year one. Does that frame the investment differently?

18 · FINAL MENTAL MODEL — PM + BUILDER MINDSET

Product Manager Mindset

- Always ask: What business problem does this solve and for whom?
- Define success metrics before writing a single line of code
- Prioritise use cases by impact times feasibility, not by technical interest
- Treat every ML model as a product: it has users, requirements, and a lifecycle
- Communicate in business language — never lead with model metrics to non-technical stakeholders

Product Builder Mindset

- Design for production from day one — notebooks are not products
- Build the smallest thing that proves value, then iterate fast
- Instrument everything: if you cannot measure it, you cannot improve it
- Automate retraining, monitoring, and deployment — they should be pipelines not manual tasks
- Document as you build — the delivery package IS the product, not the model

Deep ML Expertise + Business Outcome Focus + Fast Delivery = Premium Consulting Fees

YOUR NORTH STAR The Mindset Shift That Changes Everything

You are not selling machine learning. You are selling business outcomes enabled by machine learning. The moment you internalise this shift — in your proposals, your conversations, your deliverables — is the moment clients stop seeing you as a vendor and start seeing you as a strategic partner.