

AI Product Management

Certification Course

Exercise Workbook

by Product Faculty

Rohan Varma (OpenAI) & Henry Shi (Anthropic)

 **Note:** Do not share this workbook. Make your own copy and fill it in as you progress through each module.

All 9 Modules / From Identifying AI Opportunities to Landing the AI PM Role

Course Module Index

Module	Title
Module 1	Identify High-Leverage AI Opportunities Using the 4D Method
Module 2	Design Trustworthy AI Experiences That Drive Adoption
Module 3	Systematically Optimize Outputs for Accuracy & Relevance
Module 4	Build Rigorous Evaluation Suites — Not Vibes
Module 5	Architect Agentic AI Systems That Work
Module 6	Select the Perfect LLM for Cost, Latency & Quality
Module 7	Navigate AI Ethics, Safety & Responsible Deployment
Module 8	Lead AI Teams, Manage Stakeholders & Drive Adoption
Module 9	Build Your AI PM Portfolio & Land the Role

Module 1. Identify High-Leverage AI Opportunities Using the 4D Method

🔖 **Rohan's Rule:** Start with user pain, not technology. Ask: 'What is the most painful thing our users do today, and can AI solve it better than anything else?'

Exercise 1.1: Choose Your AI Product

Select a product you'll use throughout this entire workbook.

Field	Your Answer
Product Name	
What does it do?	
Who are the users?	
#1 User Pain	
Why could AI help here?	

Exercise 1.2: The AI Opportunity Stack

Map your product through all 4 layers. Be specific.

Layer	Key Question	Your Answer
Layer 1: User Pain	What is the most frequent, high-severity problem users face?	
Layer 2: Solution Gap	Why do existing solutions fail? What's the workaround?	
Layer 3: AI Suitability	Is this a pattern recognition problem? Is there enough data? Is the cost of being wrong acceptable?	
Layer 4: Business Value	If AI solves this, what's the measurable business impact?	

Exercise 1.3: The 4D Method — Map Your AI Opportunity

D	Step	Key Question	Your Answer
Discover	Map all user workflows	Where do users spend the most time? Where do they make the most errors?	
Define	Identify highest-pain workflow	Which workflow, if improved, would have the biggest impact?	
Determine	Apply AI Suitability Test	Is AI the right solution? Or is this a process/UX problem?	
Decide	Choose the AI approach	RAG, Agent, Classifier, or Generator — which fits?	

Exercise 1.4: 5-Question Pain Scoring Framework

$\text{Pain Score} = (\text{Sum} / 25) \times 100$

Question	What It Measures	Scoring Guide	Your Score (1-5)
Frequency	How often does this pain occur?	Daily=5 / Weekly=3 / Monthly=1	
Severity	How bad is it when it happens?	Blocks work=5 / Slows=3 / Minor=1	

Population	How many users experience this?	>80%=5 / 40-80%=3 / <40%=1	
Workaround	How painful is the current workaround?	None=5 / Bad=3 / OK=1	
Unlocks	What does solving this enable?	New revenue=5 / Retention=3 / Efficiency=1	

📌 **Your Pain Score:** Total: $__ / 25 \times 100 = __\%$ | >70 = Strong AI opportunity | 40-70 = Investigate | <40 = Probably not worth it

Exercise 1.5: Write Your 1-Page AI PRD

An AI PRD has all standard sections PLUS 7 AI-specific sections.

PRD Section	What to Include	Your Answer
Problem Statement	What is the problem? Who has it? How do we know?	
AI Approach	RAG / Agent / Classifier / Generator — and why this approach	
Data Requirements	What data does the model need? What's the quality bar?	
Model Selection	Which model tier and why (cost vs. quality trade-off)	
Prompt Strategy	High-level prompt architecture and iteration plan	
Evaluation Plan	Golden Dataset size, judge rubric, key metrics	
Guardrails Spec	Input guardrails, output guardrails, human oversight	
Cost Model	Tokens $\times$ price $\times$ volume = monthly cost estimate	
AI Risk Register	Top 3 risks with probability, severity, and mitigation	
Success Metrics	How will you know this AI feature is working?	

Module 2. Design Trustworthy AI Experiences That Drive Adoption

🔗 **Rohan's Rule:** Chat is one interaction pattern. It's not always the right one. The right pattern depends on the user's workflow, context, and cognitive load.

Exercise 2.1: Choose Your AI Interaction Pattern

Select the right pattern for your product. Justify your choice.

Pattern	Best For	Example	✓ Right for My Product?
Chat / Conversational	Open-ended exploration, Q&A, support	Customer support bot, research assistant	
Inline Suggestion	Augmenting existing workflows	GitHub Copilot, Gmail Smart Compose	
Proactive Nudge	Surfacing insights users didn't ask for	Slack summarization, CRM next-best-action	
Automated Action	Repetitive, low-stakes tasks	Auto-categorization, auto-tagging	
Structured Output	Generating documents, reports, summaries	PRD generator, meeting notes	

My chosen pattern: _____ Reason: _____

Exercise 2.2: The 3P Framework — Score Your AI Feature

P	Evaluation Question	Your Assessment	Score (1-5)
Prioritization	Is this AI feature solving the highest-priority user pain? Or is it a nice-to-have?		
Placement	Is the AI feature placed where users are already working? Or does it require them to go somewhere new?		
Prominence	Is the AI feature visible enough to be discovered? Or is it buried in a settings menu?		

Exercise 2.3: Design the 5 Trust Signals for Your Feature

Trust Signal	How to Implement It	Your Implementation Plan
Source Transparency	Show where the AI got its information.	
Confidence Indicators	Show when the AI is uncertain.	
Human Override	Always give users a way to edit, reject, or escalate AI outputs.	
Explainability	Show the reasoning, not just the answer.	
Error Recovery	Make it easy to report bad outputs and get a human alternative.	

Exercise 2.4: Design All 4 AI Error States

AI fails differently from traditional software. Design a specific UX response for each.

Error State	What Happens	Design Response (Reference)	Your UX Response
Hallucination	AI generates confident but wrong information	Source citations, confidence scores, human review flag	
Refusal	AI declines to answer (over-cautious guardrails)	Clear explanation + alternative path (human handoff)	
Irrelevance	AI answers a different question than asked	Clarification prompt, 'Did you mean...' pattern	
Timeout/Failure	Model API is slow or unavailable	Graceful degradation, cached responses, human fallback	

Module 3. Systematically Optimize Outputs for Accuracy & Relevance

📌 **Henry's Rule:** Output quality is determined by three things: the model you choose, the data you provide, and the prompt you write. PMs own two of those three.

Exercise 3.1: Prompt Engineering Hierarchy — Test All 5 Levels

Write a prompt for your product at each level. Measure the quality difference.

Level	Technique	When to Use	Quality Gain	Your Prompt (write it)
1	Zero-Shot	Simple, well-defined tasks	Baseline	
2	Few-Shot	Tasks with specific output format	+15–25%	
3	Chain-of-Thought	Reasoning tasks, multi-step problems	+20–35%	
4	Role + Context + Constraints	Domain-specific tasks	+25–40%	
5	RAG (Retrieval-Augmented)	Tasks requiring specific knowledge	+40–60%	

Exercise 3.2: Build Your Production System Prompt

A production system prompt has 6 components. Missing any one degrades quality.

Component	Template	Your Version
ROLE	"You are [specific role] at [specific context]."	
TASK	"Your job is to [specific task] given [specific inputs]."	
CONTEXT	"Here is the relevant context: [context]"	
CONSTRAINTS	"You must always [constraint]. You must never [constraint]."	
OUTPUT FORMAT	"Format your response as [format]."	
EXAMPLES	"Here are examples of good outputs: [examples]" (Few-Shot)	

Exercise 3.3: Design Your RAG Architecture

RAG solves the #1 problem with LLMs: they don't know your specific data.

RAG Lever	What It Controls	PM's Role	Your Decision
Chunking Strategy	How documents are split for retrieval	Define chunk size requirements based on use case	
Retrieval Quality	How well the search finds relevant chunks	Define relevance criteria, test with real queries	
Context Window	How much context is passed to the model	Balance quality vs. cost (more context = more tokens)	
Grounding Instructions	How strictly the model sticks to retrieved context	Define faithfulness requirements in system prompt	

Module 4. Build Rigorous Evaluation Suites — Not Vibes

📌 **Henry's Rule:** "If you can't measure it, you can't improve it. In AI, if you can't evaluate it, you're flying blind."

Exercise 4.1: Build Your Golden Dataset

Minimum viable: 30 examples. Production: 200+ examples.

Category	% of Dataset	What It Tests	Your Examples (list 2-3)
Happy Path	60%	Standard use cases your AI should handle well	
Edge Cases	25%	Unusual inputs, ambiguous queries, boundary conditions	
Adversarial	15%	Inputs designed to break your AI — jailbreaks, contradictions	

Total examples planned: _____ | Target production size: _____

Exercise 4.2: Define Your Evaluation Metrics

Set targets before you start building.

Metric	What It Measures	How to Score	Target	Your Target
Faithfulness	Is the output grounded in the provided context?	LLM-as-Judge (0-1)	>0.85	
Relevance	Does the output answer the actual question?	LLM-as-Judge (0-1)	>0.80	
Completeness	Does the output cover all required elements?	Rubric checklist	>0.75	
Latency	How fast does the system respond?	p50, p95, p99 (s)	p95 < 3s	
Cost per Query	Token cost per user interaction	\$ per 1000 queries	TBD	

Exercise 4.3: Write Your LLM-as-Judge Prompt

LLM-as-Judge scales qualitative evaluation that would otherwise require human reviewers.

📌 Template: You are an expert evaluator for an AI [use case] system. Given: User Query Retrieved Context AI Response Evaluate: Faithfulness (0-1) Relevance (0-1) Completeness (0-1) Return JSON: {faithfulness, relevance, completeness, overall, reasoning, failure_mode}

Your custom LLM-as-Judge prompt (write it below):

Exercise 4.4: Failure Pattern Analysis

After running your Golden Dataset, identify your top 3 failure patterns.

Failure Pattern	Root Cause	Frequency	Fix Plan
Failure Pattern #1			
Failure Pattern #2			
Failure Pattern #3			

Module 5. Architect Agentic AI Systems That Work

🔗 **Henry's Rule:** An AI agent is an LLM that can take actions in the world — not just generate text. It has three components: Sense (perceive inputs), Plan (decide what to do), Act (use tools to do it).

Exercise 5.1: Map Your Agent's Sense-Plan-Act Layers

Layer	Key Question	Your Answer
Sense (Inputs)	What inputs does your agent receive?	
Plan (Reasoning)	What planning pattern will you use? (ReAct / Plan-Execute / Manager-Worker)	
Act (Tools)	What tools does your agent need? List all of them.	

Exercise 5.2: Choose Your Agent Planning Pattern

Pattern	How It Works	Best For	Risk	✓ My Choice?
ReAct	Reason → Act → Observe → Repeat	Research, investigation tasks	Can loop indefinitely	
Plan-Execute	Create full plan first, then execute steps	Complex multi-step tasks	Plan may be wrong; hard to adapt	
Manager-Worker	Orchestrator agent delegates to specialist agents	Complex workflows with distinct subtasks	Coordination overhead, error propagation	

My chosen pattern: _____ Justification: _____

Exercise 5.3: Write Tool Schemas for Your Agent

Every tool your agent uses needs a schema.

Tool	Name	Description (when to use / when NOT to use)	Parameters	Returns
Tool 1				
Tool 2				
Tool 3				

Exercise 5.4: Agent Reliability Checklist

Complete this checklist before shipping any agent to production.

TERMINATION CONDITIONS

- ☐ Maximum steps defined (e.g. 15 steps)
- ☐ Maximum duration defined (e.g. 60 seconds)
- ☐ Success condition defined (how does the agent know it's done?)
- ☐ Failure condition defined (when should the agent give up?)

HUMAN OVERSIGHT

- ☐ Human approval required for irreversible actions
- ☐ Human handoff path for tasks the agent can't complete
- ☐ Audit log of all agent actions

ERROR HANDLING

- ☐ Retry logic for transient tool failures
- ☐ Graceful degradation if a tool is unavailable
- ☐ Clear error message to user if agent fails

SAFETY

- ☐ Agent cannot take actions outside defined scope
- ☐ PII is not logged in agent action history
- ☐ Rate limiting on expensive tool calls

Exercise 5.5: Agent Failure Mode Analysis

Failure Mode	What Happens	Prevention (Reference)	Your Mitigation Plan
Infinite Loop	Agent keeps calling tools without making progress	Max steps limit, loop detection	
Tool Misuse	Agent calls the wrong tool or with wrong parameters	Clear tool descriptions, parameter validation	
Context Loss	Agent forgets earlier steps in long tasks	Explicit state object, summarization at checkpoints	
Hallucinated Actions	Agent invents tool calls that don't exist	Strict tool list in system prompt, output validation	
Scope Creep	Agent takes actions beyond its intended scope	Explicit action constraints, human approval	

Module 6. Select the Perfect LLM for Cost, Latency & Quality

Henry's Rule: Model selection is one of the highest-leverage decisions an AI PM makes. The wrong model costs you 10x more than necessary, or delivers 30% worse quality.

Exercise 6.1: 3-Dimension Model Selection Matrix

Score your AI feature on all three dimensions before selecting a model.

Dimension	Questions to Ask	Drives Toward	Your Answer
Quality Requirements	What's the minimum acceptable quality? What's the cost of a wrong answer?	Higher quality → larger model	
Latency Requirements	What's the maximum acceptable response time? Is this real-time or async?	Lower latency → smaller/faster	
Cost Constraints	What's the budget per query? What's the expected query volume?	Lower cost → smaller or caching	

Exercise 6.2: Select Your Model Tier

Tier	Models	Best For	Cost	Latency	✓ My Choice?
Frontier	GPT-4o, Claude 3.5 Sonnet, Gemini 1.5 Pro	Complex reasoning, high-stakes outputs	High (\$\$\$)	Medium (2–5s)	
Mid-Tier	GPT-4o-mini, Claude 3 Haiku, Gemini Flash	Most production use cases	Medium (\$)	Fast (0.5–2s)	
Small/Fast	GPT-3.5-turbo, Claude Instant	Simple classification, routing	Low (\$)	Very fast (<0.5s)	
Open Source	Llama 3, Mistral, Phi-3	Privacy-sensitive, on-premise	Infra cost only	Depends on infra	

Exercise 6.3: Build Your Cost Model

Calculate monthly cost at 3 volume scenarios.

Variable	Scenario A	Scenario B	Scenario C
Avg Input Tokens / Query			
Avg Output Tokens / Query			
Input Price (per 1M tokens)			
Output Price (per 1M tokens)			
Cost per Query (\$)			
Monthly Query Volume	10,000	100,000	1,000,000
MONTHLY COST (\$)			

Exercise 6.4: Fine-Tuning vs. Prompt Engineering vs. RAG Decision

Approach	When to Use	Cost	Time to Value	✓ My Choice?
Prompt Engineering	Always try first. Solves 80% of quality problems.	Low	Hours	
RAG	When the model needs your specific data/knowledge.	Medium	Days–Weeks	
Fine-Tuning	When you need specific style/format/domain expertise that prompting can't achieve.	High	Weeks–Months	

My decision: _____ Reasoning: _____

Module 7. Navigate AI Ethics, Safety & Responsible Deployment

📌 **Henry's Rule:** "At Anthropic, every PM is responsible for the ethical dimensions of their product. Ethics as product quality means: ethical requirements are defined in the PRD, guardrails are native to the UX — not bolted on."

Exercise 7.1: Harm Assessment

Identify the top 3 harms your AI feature could cause. Score probability × severity.

#	Potential Harm	Probability (VL/L/M/H)	Severity (Min/Mod/Sev/Crit)	Risk Score (P×S)
Harm 1				
Harm 2				
Harm 3				

Exercise 7.2: AI Privacy Checklist

Map every data flow. Check every box before launch.

DATA COLLECTION

- ☐ What user data does the AI receive as input?
- ☐ Is any of this data PII? (names, emails, health info, location)
- ☐ Does the user know their data is being processed by AI?
- ☐ Have users consented to AI processing of their data?

DATA STORAGE

- ☐ Are conversation logs stored? For how long?
- ☐ Are conversation logs used for model training?
- ☐ Can users request deletion of their conversation history?

DATA TRANSMISSION

- ☐ Is data encrypted in transit to the model provider?
- ☐ Does your model provider's DPA meet compliance requirements?
- ☐ Is any data sent to third-party tools or APIs?

Exercise 7.3: Design Your 3 Lines of Defense

Line of Defense	What It Covers (Reference)	Your Implementation Plan
Line 1: Input Guardrails	Content classification, PII detection, prompt injection detection, rate limiting	
Line 2: Model-Level Safety	System prompt constraints, negative instructions, safe messaging guidelines	
Line 3: Output Guardrails	Content moderation, PII detection in outputs, factual claim flagging, human review	

Exercise 7.4: EU AI Act Classification

Risk Level	Examples	Requirements	My Feature Falls Here?
Unacceptable (Banned)	Social scoring, real-time biometric surveillance	Do not build. Full stop.	

High Risk	AI in hiring, credit scoring, medical devices	Mandatory human oversight, transparency, accuracy requirements	
Limited Risk	Chatbots, AI-generated content	Disclose AI interaction, label AI-generated content	
Minimal Risk	Spam filters, recommendation systems	No specific requirements	

Exercise 7.5: Pre-Launch Ethics Review

Complete this review and get sign-off from one non-PM stakeholder before launch.

Section	Key Questions	Your Answers
Harm Assessment	Worst-case harm? Probability? Severity? Risk Score?	
Fairness Assessment	Which user groups? Tested across all groups? Disproportionate harm?	
Transparency	Do users know they're interacting with AI? Human appeal mechanism?	
Guardrails Assessment	Input guardrails? Output guardrails? Human oversight mechanism?	
Incident Response	What is the plan if this AI feature causes harm?	

Launch Decision: ☐APPROVED ☐APPROVED WITH CONDITIONS ☐DELAYED ☐REJECTED

Signatures: PM: _____ Legal: _____ Trust & Safety: _____

Module 8. Lead AI Teams, Manage Stakeholders & Drive Adoption

 **Henry's Rule:** Always anchor AI performance to a human or legacy system baseline. Never present AI metrics in isolation.

Exercise 8.1: Translate Technical Metrics to Business Language

Use this framework to communicate AI performance to non-technical stakeholders.

Technical Metric	Business Translation (Reference)	Your Translation
Faithfulness: 0.87	13% of responses contain information not in our knowledge base	
Task Completion Rate: 78%	22% of users don't complete their intended task — 220 users/day needing human support	
p95 Latency: 4.2s	1 in 20 users waits more than 4 seconds — target is under 3 seconds	
Your Metric: __	Your translation:	Your translation:

Exercise 8.2: Map Your Stakeholders — 4 Archetypes

Archetype	Profile	How to Win Them (Reference)	Your 3-Sentence Plan
The Skeptic	"AI is just hype. Show me the ROI."	Lead with baseline comparison. Hard numbers. Working prototype with real metrics.	
The Enthusiast	"AI can do everything!"	Channel their energy. Show 90-day roadmap + bigger vision.	
The Risk Manager	"What if it goes wrong?"	Show your risk register. 5 failure modes + mitigation for each.	
The End User Champion	"Will this actually help users?"	Involve them in user research. Show real user feedback from pilots.	

Exercise 8.3: Build Your AI Roadmap — Horizon Model

$Priority\ Score = (Business\ Impact \times 0.30) + (Feasibility \times 0.25) + (User\ Pain \times 0.25) + (Strategic\ Value \times 0.20)$

Feature	Business Impact (1-10)	Feasibility (1-10)	User Pain (1-10)	Strategic Value (1-10)	Priority Score	Horizon (Now/Next/Later)
Feature 1						
Feature 2						
Feature 3						
Feature 4						
Feature 5						

Exercise 8.4: Design Your AI Adoption Strategy

Identify your top 2 adoption barriers and build a mitigation plan.

Barrier	How to Overcome It (Reference)	Your Mitigation Plan
Trust Deficit	Show sources, confidence levels, start with low-stakes use cases	
Workflow Disruption	Meet users where they are, make AI opt-in, show time savings	
Black Box Problem	Explain reasoning, show inputs, make uncertainty visible	
Fear of Replacement	Frame as augmentation, involve users in design	
Learning Curve	Provide prompt templates, show examples, progressive disclosure	

Exercise 8.5: Define Your Adoption Health Score

$Adoption\ Health\ Score = (Acceptance\ Rate \times 0.30) + (Return\ Rate \times 0.25) + (Task\ Completion \times 0.25) + (User\ Satisfaction \times 0.20)$

Metric	What It Measures	Target (Reference)	Your Target
AI Acceptance Rate	% of AI suggestions users accept without editing	>60%	
AI Override Rate	% of AI outputs users significantly edit or reject	<30%	
AI Return Rate	% of users who use AI feature again after first use	>50%	
Time-to-Value	Time from first AI use to first 'aha moment'	<5 min	
Task Completion Rate	% of users who complete their intended task with AI help	>75%	

Module 9. Build Your AI PM Portfolio & Land the Role

📌 **Rohan's Rule:** Demand for AI PMs is growing 40% year-over-year. Most PMs have surface-level AI knowledge. Genuinely qualified AI PMs are scarce. This gap is your opportunity.

Exercise 9.1: Identify Your AI PM Profile

Which profile best describes you? Build your portfolio strategy around it.

Profile	Description	Key Differentiator	✓ My Profile?
The Builder	Has shipped AI products. Has a portfolio. Can talk about failures as fluently as successes.	Commands a premium	
The Translator	Deep domain expertise (healthcare, finance, legal) + AI knowledge.	Identifies AI opportunities generalists miss	
The Evaluator	Obsessed with quality and measurement. Has built eval suites. Understands LLM-as-Judge.	Rare and highly valued	

Exercise 9.2: Build Your 5-Artifact AI PM Portfolio

Check off each artifact as you complete it. Every artifact should be publicly shareable.

#	Artifact	Must Include	Status (Draft/Done/Published)
Artifact 1	AI Opportunity Assessment (2 pages)	Problem, AI Opportunity, Business Case, Risk Register	
Artifact 2	The AI PRD (4–6 pages)	AI Approach, Data Req, Model Selection, Prompt Strategy, Eval Plan, Guardrails, Cost Model, Risk Register	
Artifact 3	The Eval Suite	30-example Golden Dataset, custom judge prompt, scoring spreadsheet, failure analysis, improvement log	
Artifact 4	Agent Architecture Document	Agent Overview, Sense/Plan/Act layers, Tool Schemas, State Object, Termination Conditions, Failure Mode Analysis	
Artifact 5	STAR-AI Case Study	Situation, Task, Action (AI-specific), Result — with eval scores, cost model, failure rate, next steps	

Exercise 9.3: Write Your STAR-AI Case Study

Use this template. Be specific. Include numbers. Talk about failures.

STAR Component	What to Include	Your Answer
SITUATION	Business context, problem, baseline (before AI)	
TASK	Your specific role, responsibilities, constraints	
ACTION	AI approach chosen and why Eval criteria Prompt iterations Failure modes encountered Responsible AI considerations	
RESULT	Key metric change Cost/latency impact Learnings What you'd do differently	
AI ADDITIONS	Eval scores before/after Cost model at scale Failure rate Next steps	

Exercise 9.4: AI PM Interview Prep — 5 Question Types

Practice answering each type. Time yourself. Aim for 2.5 minutes per answer.

Question Type	Example Question	Answer Structure (Reference)	Your Draft Answer
Type 1: AI Opportunity	"How would you use AI to improve [product]?"	(1) Highest-pain workflow (2) AI Suitability Test (3) AI approach (4) Success metrics (5) Top risk	
Type 2: Technical Depth	"Walk me through how RAG works."	(1) Plain English (2) When to use vs. alternatives (3) 4 quality levers (4)	

		When RAG fails (5) Portfolio example	
Type 3: Metrics	"How would you measure success of an AI feature?"	4-Layer KPI: Model → Agent → Product → Business	
Type 4: Failure	"Tell me about an AI project that didn't go as planned."	(1) What you built (2) What happened (3) Root cause (4) What you changed (5) What you'd do differently	
Type 5: Responsible AI	"How do you approach AI ethics?"	(1) Ethics as product quality (2) 6 dimensions (3) Pre-launch review (4) Specific trade-off (5) Your philosophy	

Exercise 9.5: Your 3-Platform Portfolio Strategy

Publish your artifacts publicly. Your portfolio is your proof of work.

Platform	What to Publish	Goal	Your Link / Status
LinkedIn	1 post per artifact (5 total), learnings not just outputs, #AIProductManagement	Visibility and inbound opportunities	
GitHub	Eval suite code, Golden Dataset (anonymized), prompt templates	Technical credibility signal	
Personal Site/Notion	All 5 artifacts linked, AI PM Philosophy statement, contact info	Single URL to share with recruiters	

✓ Course Completion Checklist

🔗 **You're ready to land the AI PM role when:** You can check every box below AND explain the reasoning behind each decision to a skeptical stakeholder.

MODULE 1 — AI Opportunity Identification

- ☐ Completed AI Opportunity Stack for my product
- ☐ Scored my opportunity using the 5-Question Framework (Pain Score > 70)
- ☐ Wrote a 1-page AI PRD with all 9 AI-specific sections

MODULE 2 — AI UX Design

- ☐ Chose the right interaction pattern and justified it
- ☐ Scored my feature on the 3P Framework
- ☐ Designed all 5 trust signals
- ☐ Designed all 4 AI error states

MODULE 3 — Output Optimization

- ☐ Tested prompts at all 5 levels of the hierarchy
- ☐ Built a production system prompt with all 6 components
- ☐ Designed my RAG architecture with all 4 quality levers

MODULE 4 — Evaluation Suites

- ☐ Built a 30-example Golden Dataset (60/25/15 split)
- ☐ Defined evaluation metrics with specific targets
- ☐ Written a custom LLM-as-Judge prompt
- ☐ Identified top 3 failure patterns and root causes

MODULE 5 — Agentic AI

- ☐ Mapped Sense-Plan-Act layers
- ☐ Chose a planning pattern and justified it
- ☐ Written tool schemas for all agent tools
- ☐ Completed the Agent Reliability Checklist

MODULE 6 — LLM Selection & Cost

- ☐ Scored my feature on the 3-dimension matrix
- ☐ Selected a model tier with justification
- ☐ Built cost model at 3 volume scenarios
- ☐ Decided: Prompt Engineering / RAG / Fine-Tuning

MODULE 7 — AI Ethics & Safety

- ☐ Completed Harm Assessment (top 3 harms scored)
- ☐ Completed AI Privacy Checklist
- ☐ Designed 3 Lines of Defense
- ☐ Classified feature under EU AI Act
- ☐ Completed Pre-Launch Ethics Review with sign-off

MODULE 8 — Leadership & Adoption

- ☐ Translated technical metrics to business language
- ☐ Mapped stakeholders using 4 archetypes
- ☐ Built AI Roadmap using Horizon Model
- ☐ Defined Adoption Health Score with targets

MODULE 9 — Portfolio & Job Search

- ☐ Identified my AI PM profile (Builder/Translator/Evaluator)
- ☐ Completed all 5 portfolio artifacts
- ☐ Written STAR-AI case study
- ☐ Practiced all 5 interview question types
- ☐ Published portfolio on LinkedIn + GitHub + Personal Site

Good luck. Now go build something real. 🚀